


Aggregating more frames of one static camera



5 % target

175 fixed-camera views; statistical unit = view. Every point is recomputed from ONE inference run (CAM-EXP-004.1), so no step in the curve is a run boundary. For $N \leq 8$ all $C(8, N)$ frame subsets are evaluated and averaged per view. The crosses show the CAM-EXP-004 original run at the same N for comparison; they differ by at most 0.52 pp, which is the GeoCalib run-to-run variation, and are deliberately not joined into the curve.